

Ricardo Perez Castillo

ricardoperez@berkeley.edu LinkedIn: ricardoperezcastillo +1 (510) 365-6415
GitHub: ricardo-pc Portfolio: ricardo-pc.github.io Berkeley, CA

Summary

Data scientist with 4+ years at McKinsey and Deloitte building ML models, data pipelines, and analytics products for enterprise clients across the U.S., Europe, and Latin America. Completed an M.A. in Statistics at UC Berkeley in May 2026, with coursework and project work in statistical inference, experimentation, and machine learning. Comfortable owning problems end-to-end: from scoping and pipeline work to modeling, stakeholder communication, and production deployment.

Education

University of California, Berkeley Aug 2025 – May 2026
M.A. in Statistics (completed May 2026) Berkeley, CA
Coursework: Statistical Inference, Statistical Computing, Linear Models, Machine Learning, Bayesian Statistics.

Tecnológico de Monterrey, Campus Monterrey Jan 2015 – Dec 2019
B.S. in Chemical Engineering Monterrey, México
Cumulative GPA: 91.38/100 Class Rank: 3/19

Experience

McKinsey & Company Jul 2022 – Jul 2025
Senior Software Delivery Analyst, Data (Senior Data Engineer/Analyst) Jul 2024 – Jul 2025
Software Delivery Analyst, Data (Data Engineer/Analyst) Jul 2022 – Jul 2024
San Pedro Garza García, México

- Owned analytics and ML deployment end-to-end for Spendscope across U.S., Germany, and Latin America, partnering cross-functionally to define success metrics, ship data features, and surface insights that drove user adoption and client retention.
- Built a Snowflake data warehouse consolidating 16 ERP systems (SAP, Oracle, Coupa) using dimensional modeling (star schema with fact and dimension tables); automated ingestion (REST APIs, SFTP → AWS S3 → Python loaders) cut refresh time from 2–3 days to 24 hours, with a single extraction reaching 400M rows of transactional data.
- Quantified a \$340–\$460M raw-material price arbitrage opportunity for a U.S. tire manufacturer through forecasting and causal analysis on spend data, presenting findings directly to C-suite stakeholders and informing procurement strategy.
- Shipped an end-to-end production NLP text classification pipeline for spend categorization, writing model-generated labels directly to warehouse records and cutting manual classification from weeks to days.
- *Equal at McKinsey:* Board member and Monterrey office leader for McKinsey’s LGBT+ network; organized office-wide D&I initiatives aligned with the firm’s strategic vision.

Deloitte Jan 2021 – Jun 2022
Consulting Analyst Monterrey, México

- Reduced supplier harmonization time by 60% (5 to 2 days) by building and deploying a spaCy NLP entity-resolution model trained on accounts payable and purchase-order data, processing millions of rows across U.S. client engagements.
- Improved ML categorization accuracy by 15 percentage points by retraining and redeploying a Google Cloud-hosted classification model, while maintaining ETL pipelines in Alteryx and Python across AP, PO, T&E, and P-card data sources.
- Analyzed \$114B in combined spend across multiple industries, building Tableau dashboards that became the primary decision-making interface for finance and procurement stakeholders.

Projects

Two-Stage Recommendation System: A Sparsity Analysis UC Berkeley, CS 289A · Spring 2026
Python, PyTorch, Bayesian Optimization, BPR, Matrix Factorization, NCF

- Trained and evaluated three recommendation systems (MF, NCF, two-stage BPR ranker) across five data density levels on MovieLens 1M (1M ratings, 6K users), using Bayesian optimization to independently tune each model; found that Bayesian-optimized MF outperformed NCF at every density level on NDCG@10, with the BPR reranker improving HR@10 by up to 1.2 points in sparse regimes.

Linguistic Predictors of Click-Through Rate in Headline A/B Tests UC Berkeley, STAT 230A · Spring 2026

Python, OLS, Cluster-Robust Inference, Feature Engineering, VADER Sentiment, A/B Testing

- Fit a cluster-robust OLS regression on 54,956 headline packages across 11,649 A/B tests from the Upworthy Research Archive; found that exclamation marks reduce CTR by 25% and numeric openings increase it by 6%, with results stable across HC3, cluster-robust, and impression-weighted specifications.

PECARN Traumatic Brain Injury Analysis UC Berkeley, STAT 214 · Mar 2026

Python, Bayesian Inference, Classification, scikit-learn

- Applied EDA, Bayesian inference, and classification models to 43,399 pediatric head trauma patients to validate and extend the Kuppermann clinical decision rule for CT triage.

Berkeley Grocery Tracker OpenAI Codex Challenge · Apr 2026

Python, Data Engineering, Web App, Claude Code / OpenAI Codex

- Built an end-to-end live price tracker for Trader Joe's, Whole Foods, and Safeway across Berkeley using AI-assisted development with Claude Code and OpenAI Codex, from data pipeline design to web app deployment.

Digital Water: Predicting Monterrey's Water Crisis Atos IT Challenge · 2020

Python, Time Series Forecasting, Machine Learning

- Built an ML forecasting system that predicted Monterrey's water supply crisis two years before 5 million people faced severe rationing, surfacing early warning signals from historical reservoir and consumption data.

Skills

Analytics & ML: A/B Testing, Experimentation Design, Statistical Inference, Causal Inference, Forecasting, Regression, Classification, Clustering, NLP, Recommendation Systems, Feature Engineering

Programming & Data: Python (pandas, NumPy, scikit-learn, statsmodels, spaCy), R, SQL, Snowflake, Microsoft SQL Server, AWS S3, ETL, Dimensional Modeling

BI & Infrastructure: Tableau, Power BI, Azure, GCP, Git, REST APIs, SnapLogic